

# TITANIC - Machine Learning from Disaster



*N M A Chamindu Nipun*

# Titanic - Machine Learning from Disaster

## INTRODUCTION



Titanic disaster is one of the most famous shipwrecks in the world history.

Titanic was a British cruise liner that sank in the North Atlantic Ocean a few hours after colliding with an iceberg.

While there are facts available to support the cause of the shipwreck, there are various speculations regarding the survival rate of passengers in the Titanic disaster.

Over the years, data of survived as well as deceased passengers has been collected.

The dataset is publically available on a website called Kaggle.com.

This dataset has been studied and analyzed using Logistic Regression machine learning algorithm.

Pre-Processing

Model Training

RESULTS

CONCLUSION

# Description of Data

## INTRODUCTION

- **Code Link:**

[https://github.com/Chamindu77/TITANIC-Machine\\_Learning\\_from\\_Disaster/blob/main/titanic\\_ptoject.ipynb](https://github.com/Chamindu77/TITANIC-Machine_Learning_from_Disaster/blob/main/titanic_ptoject.ipynb)

- **Dataset:** Titanic - Machine Learning from Disaster

- **No. Data Points :** 891

- **No. Rows & Columns:**

```
titanic_df.shape  
  
(891, 12)
```

Pre-Processing

Model Training

RESULTS

CONCLUSION

# Description of Data

## INTRODUCTION

### ○ Data Dictionary :

Variable	Definition	Key
survival	Survival	0 = No, 1 = Yes
pclass	Ticket class	1 = 1st, 2 = 2nd, 3 = 3rd
sex	Sex	
Age	Age in years	
sibsp	# of siblings / spouses aboard the Titanic	
parch	# of parents / children aboard the Titanic	
ticket	Ticket number	
fare	Passenger fare	
cabin	Cabin number	
embarked	Port of Embarkation	C = Cherbourg, Q = Queenstown, S = Southampton

Pre-Processing

Model Training

RESULTS

CONCLUSION

# Pre-Processing Data

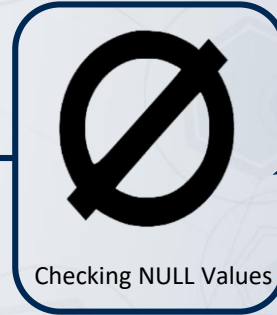
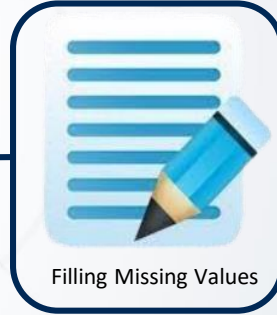



INTRODUCTION

**Pre-Processing**

Model Training

RESULTS

CONCLUSION



# Pre-processing Data

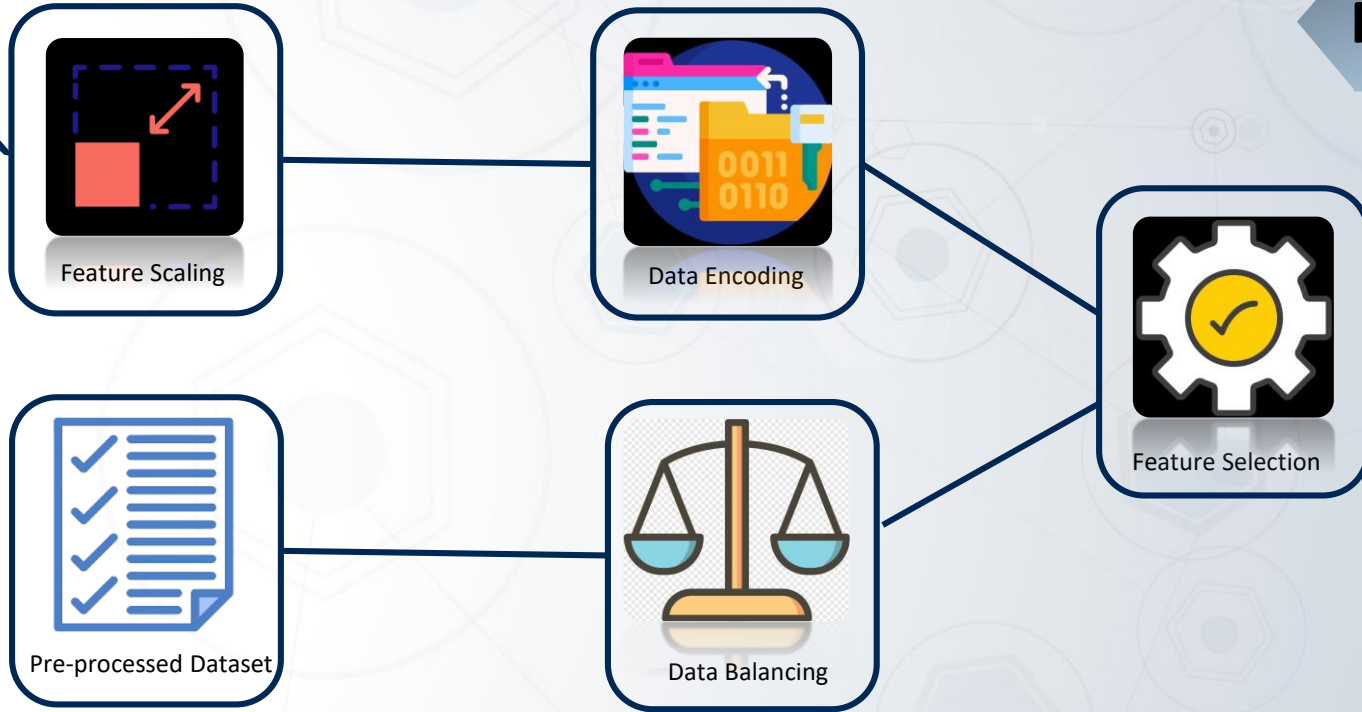
INTRODUCTION

## Pre-Processing

Model Training

RESULTS

CONCLUSION





# Describe Dataset

INTRODUCTION

Pre-Processing

Model Training

RESULTS

CONCLUSION

```
titanic_df.head()
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S





# Describe Dataset

INTRODUCTION

Pre-Processing

Model Training

RESULTS

CONCLUSION

```
titanic_df.describe()
```

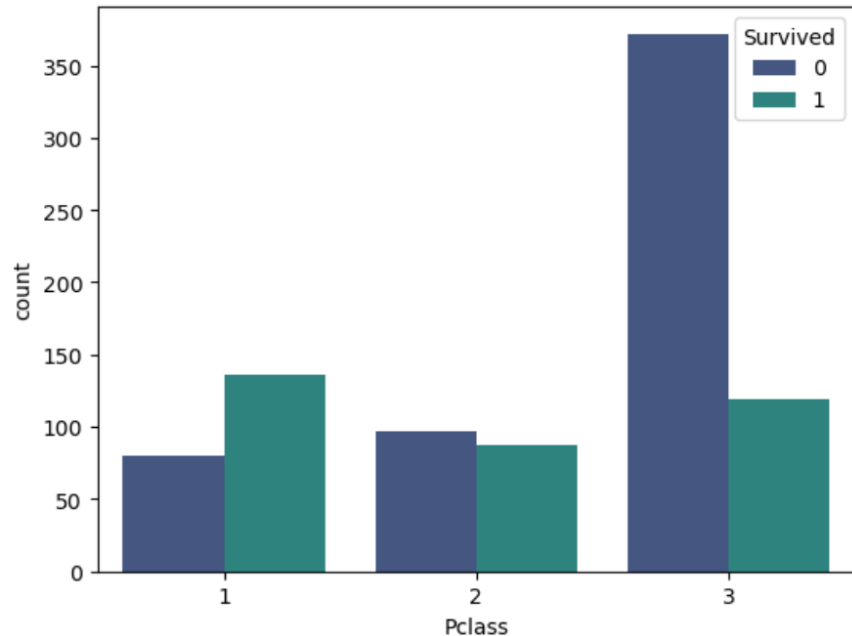
	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
<b>count</b>	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
<b>mean</b>	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
<b>std</b>	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
<b>min</b>	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
<b>25%</b>	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
<b>50%</b>	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
<b>75%</b>	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
<b>max</b>	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200



# Survival vs Ticket class

```
custom_palette = sns.color_palette("viridis", 3)
sns.countplot(x='Pclass', hue='Survived', data=titanic_df_3, palette=custom_palette)

# Group DataFrame by 'Pclass' and calculate mean survival rate for each class
titanic_df_3.groupby(['Pclass'], as_index=False)['Survived'].mean()
```



Pclass	Survived
1	0.629630
2	0.472826
3	0.242363

INTRODUCTION

Pre-Processing

Model Training

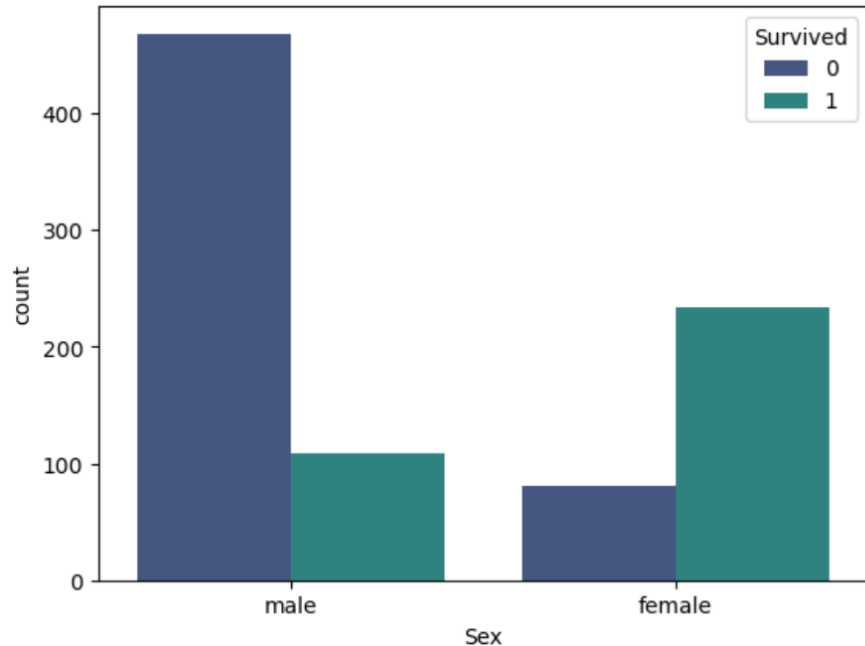
RESULTS

CONCLUSION

# Survival vs Gender

```
custom_palette = sns.color_palette("viridis", 3)
sns.countplot(x='Sex', hue='Survived', data=titanic_df_6, palette=custom_palette)

# Group the DataFrame by the 'Sex' column and calculate the mean survival rate for gender
titanic_df_6.groupby(['Sex'], as_index=False)['Survived'].mean()
```



Sex	Survived
female	0.742038
male	0.188908

# Survival vs Age in years

INTRODUCTION

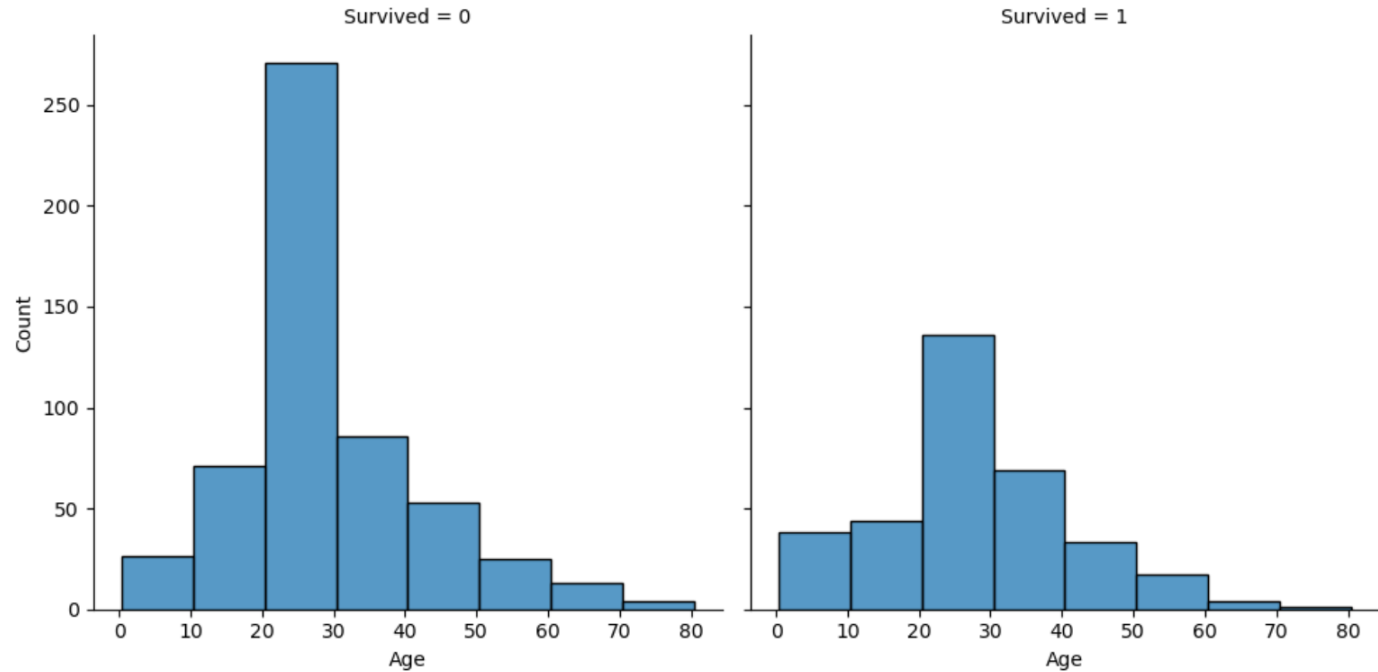
```
sns.displot(titanic_df_6, x='Age', col='Survived', binwidth=10, height=5)
```

Pre-Processing

Model Training

RESULTS

CONCLUSION



# Survival vs siblings

INTRODUCTION

```
custom_palette = sns.color_palette("viridis", 3)
sns.countplot(x='SibSp', hue='Survived', data=titanic_df_6, palette=custom_palette)
```


```
# Group the DataFrame by the 'SibSp' column and calculate the mean survival rate for each unique value
titanic_df_6.groupby(['SibSp'], as_index=False)['Survived'].mean()
```

## Pre-Processing

Model Training

RESULTS

CONCLUSION



SibSp	Survived
0	0.345395
1	0.535885
2	0.464286
3	0.250000
4	0.166667
5	0.000000
8	0.000000

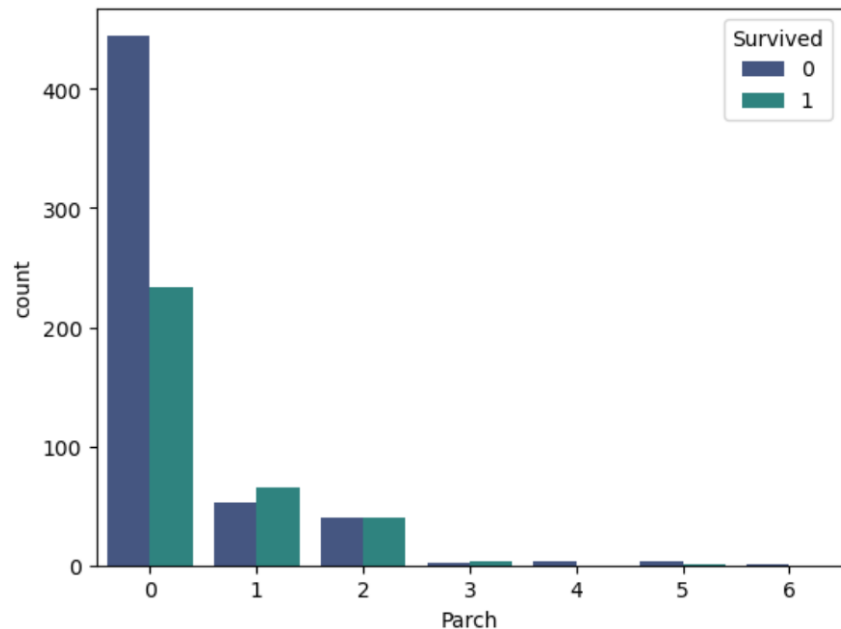
# Survival vs Parents

INTRODUCTION

```
custom_palette = sns.color_palette("viridis", 3)
sns.countplot(x='Parch', hue='Survived', data=titanic_df_6, palette=custom_palette)
```

```
# Group the DataFrame by the 'Parch' column and calculate the mean survival rate for each unique value
titanic_df_6.groupby(['Parch'], as_index=False)['Survived'].mean()
```

Pre-Processing



	Parch	Survived
0	0	0.343658
1	1	0.550847
2	2	0.500000
3	3	0.600000
4	4	0.000000
5	5	0.200000
6	6	0.000000

Model Training

RESULTS

CONCLUSION

# Survival vs Passenger fare

INTRODUCTION


```
sns.displot(titanic_df_9, x='Fare', col='Survived', binwidth=80, height=5)
```

## Pre-Processing

Model Training

RESULTS

CONCLUSION

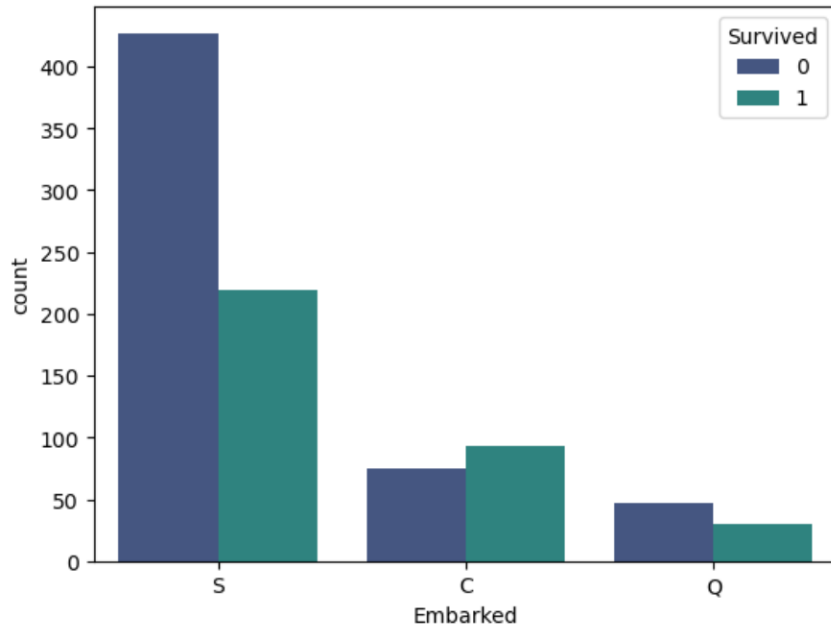


# Survival vs Port of Embarkation

INTRODUCTION

```
# Calculate the mean survival rate for passengers grouped by their port of embarkation  
titanic_df_10.groupby(['Embarked'], as_index=False)['Survived'].mean()
```

```
custom_palette = sns.color_palette("viridis", 3)  
sns.countplot(x='Embarked', hue='Survived', data=titanic_df_10, palette=custom_palette)
```



Embarked	Survived
----------	----------

C	0.553571
---	----------

Q	0.389610
---	----------

S	0.339009
---	----------

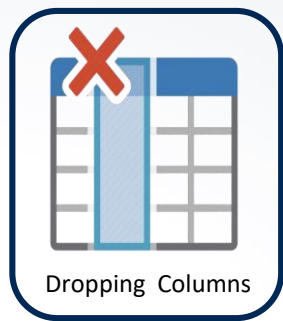
Pre-Processing

Model Training

RESULTS

CONCLUSION





# Dropping Columns

INTRODUCTION

Pre-Processing

Model Training

RESULTS

CONCLUSION

## ❑ PassengerId Column

```
# Drop the 'PassengerId' column from the DataFrame 'titanic_df_3'  
titanic_df_3 = titanic_df_3.drop(columns = 'PassengerId',axis=1)
```

## ❑ Ticket Column

```
# drop Ticket columns  
titanic_df_9 = titanic_df_9.drop(columns = 'Ticket',axis=1)
```

## ❑ Cabin Column

```
# drop Cabin columns  
titanic_df_10 = titanic_df_10.drop(columns = 'Cabin',axis=1)
```

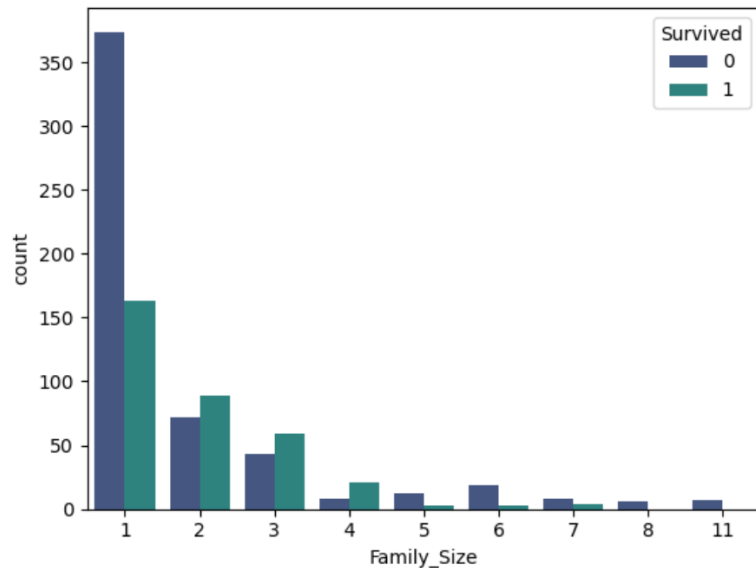
# Create Columns – Family\_Size

INTRODUCTION

```
# create family size column using SibSp column and Parch column
titanic_df_7['Family_Size'] = titanic_df_7['SibSp'] + titanic_df_7['Parch'] + 1

# Calculate the mean survival rate for passengers grouped by their family size
titanic_df_7.groupby(['Family_Size'], as_index=False)['Survived'].mean()
```

Pre-Processing



Family_Size	Survived
1	0.303538
2	0.552795
3	0.578431
4	0.724138
5	0.200000
6	0.136364
7	0.333333
8	0.000000
11	0.000000

Model Training

RESULTS

CONCLUSION

# Create Columns – Title

INTRODUCTION

Pre-Processing

Model Training

RESULTS

CONCLUSION

```
titanic_df_3['Name']
```

```
0      Braund, Mr. Owen Harris
1  Cumings, Mrs. John Bradley (Florence Briggs Th...
2      Heikkinen, Miss. Laina
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)
4      Allen, Mr. William Henry
...
886  Montvila, Rev. Juozas
887  Graham, Miss. Margaret Edith
888  Johnston, Miss. Catherine Helen "Carrie"
889      Behr, Mr. Karl Howell
890      Dooley, Mr. Patrick
Name: Name, Length: 891, dtype: object
```

```
titanic_df_4['Title']
```

```
0      Mr
1      Mrs
2      Miss
3      Mrs
4      Mr
...
886    Rev
887    Miss
888    Miss
889     Mr
890     Mr
Name: Title, Length: 891, dtype: object
```

# Create Columns – Title

INTRODUCTION

## Pre-Processing

Model Training

RESULTS

CONCLUSION

	Title
0	Capt
1	Col
2	Don
3	Dr
4	Jonkheer
5	Lady
6	Major
7	Master
8	Miss
9	Mlle
10	Mme
11	Mr
12	Mrs
13	Ms
14	Rev
15	Sir
16	the Countess

```
# Map titles to corresponding categories:

#military - Capt, Col, Major
#noble - Jonkheer, the Countess, Don, Lady, Sir
#unmarried Female - Mlle, Ms, Mme

# Replace titles with corresponding categories
titanic_df_5 = titanic_df_4.copy()
titanic_df_5['Title'] = titanic_df_5['Title'].replace({
    'Capt': 'Military',
    'Col': 'Military',
    'Major': 'Military',
    'Jonkheer': 'Noble',
    'the Countess': 'Noble',
    'Don': 'Noble',
    'Lady': 'Noble',
    'Sir': 'Noble',
    'Mlle': 'Noble',
    'Ms': 'Noble',
    'Mme': 'Noble'
})
```

	Title	count
0	Dr	7
1	Master	40
2	Military	5
3	Miss	182
4	Mr	517
5	Mrs	125
6	Noble	9
7	Rev	6

# Create Columns – Title

INTRODUCTION

## Pre-Processing

Model Training

RESULTS

CONCLUSION

	Title
0	Capt
1	Col
2	Don
3	Dr
4	Jonkheer
5	Lady
6	Major
7	Master
8	Miss
9	Mlle
10	Mme
11	Mr
12	Mrs
13	Ms
14	Rev
15	Sir
16	the Countess

```
# Map titles to corresponding categories:

#military - Capt, Col, Major
#noble - Jonkheer, the Countess, Don, Lady, Sir
#unmarried Female - Mlle, Ms, Mme

# Replace titles with corresponding categories
titanic_df_5 = titanic_df_4.copy()
titanic_df_5['Title'] = titanic_df_5['Title'].replace({
    'Capt': 'Military',
    'Col': 'Military',
    'Major': 'Military',
    'Jonkheer': 'Noble',
    'the Countess': 'Noble',
    'Don': 'Noble',
    'Lady': 'Noble',
    'Sir': 'Noble',
    'Mlle': 'Noble',
    'Ms': 'Noble',
    'Mme': 'Noble'
})
```

	Title	count
0	Dr	7
1	Master	40
2	Military	5
3	Miss	182
4	Mr	517
5	Mrs	125
6	Noble	9
7	Rev	6

# Checking NULL Values



INTRODUCTION

Pre-Processing

Model Training

RESULTS

CONCLUSION



```
titanic_df.isnull().sum()
```


PassengerId	0
Survived	0
Pclass	0
Name	0
Sex	0
Age	177
SibSp	0
Parch	0
Ticket	0
Fare	0
Cabin	687
Embarked	2
dtype:	int64

# Filling Missing Values

INTRODUCTION

## ❑ Age column

```
# Fill missing values in the 'Age' column with the mean age of all passengers  
titanic_df_1['Age'].fillna(titanic_df_1['Age'].mean(), inplace=True)
```



Pre-Processing

Model Training

RESULTS

CONCLUSION

## ❑ Embarkation column

```
# Fill missing values in the 'Embarked' column with the mode of the column  
titanic_df_2['Embarked'].fillna(titanic_df_2['Embarked'].mode()[0], inplace=True)
```



# Outlier Detection

- ☐ Age Column
- ☐ Fare Column
- ☐ Family\_Size Column

$$z = \frac{x - \mu}{\sigma}$$

$$\sigma = \sqrt{\frac{\sum (x_i - \mu)^2}{N}}$$

$$\mu = \frac{\sum x}{N}$$

INTRODUCTION

**Pre-Processing**

Model Training

RESULTS

CONCLUSION

# Outlier Detection

INTRODUCTION

## □ Age Column

### Step - 01

```
age_mean = np.mean(titanic_df_10['Age'])  
age_mean
```

```
29.69911764705882
```

```
age_std = np.std(titanic_df_10['Age'])  
age_std
```

```
12.994716872789033
```

### Step - 02

```
titanic_df_11 = titanic_df_10.copy()  
titanic_df_11['Age_z_score'] = (titanic_df_11['Age'] - age_mean)/age_std
```

### Step - 03

```
titanic_df_11['Age_z_score'].min()
```

```
-2.2531554887793948
```

```
titanic_df_11['Age_z_score'].max()
```

```
3.8708717431367314
```

Pre-Processing

Model Training

RESULTS

CONCLUSION

# Outlier Detection

INTRODUCTION

## Step - 04

```
age_outlier_indexes = []  
age_outlier_indexes.extend(titanic_df_11.index[titanic_df_11['Age_z_score'] > 3].tolist())  
age_outlier_indexes
```

[96, 116, 493, 630, 672, 745, 851]



Threshold Value

## Step - 05

```
titanic_df_11_new = titanic_df_11.drop(titanic_df_11.index[age_outlier_indexes])
```

## Step - 06

```
titanic_df_12 = titanic_df_11_new.drop('Age_z_score', axis=1)
```



Pre-Processing

Model Training

RESULTS

CONCLUSION

# Feature Scaling

INTRODUCTION

Pre-Processing

Model Training

RESULTS

CONCLUSION

Normalization




$$x' = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$$

Standardization



$$x' = \frac{x - \mu}{\sigma}$$



# Feature Scaling

❑ Age Column

❑ Fare Column

INTRODUCTION

Pre-Processing

Model Training

RESULTS

CONCLUSION

```
titanic_df_17 =titanic_df_16.copy()

scaler = MinMaxScaler(feature_range=(0, 10))

numeric_columns = ['Age', 'Fare']

for column in numeric_columns:
    column_data = titanic_df_17[column].values.reshape(-1, 1)
    scaled_data = scaler.fit_transform(column_data)
    titanic_df_17[column] = scaled_data
```

# Data Encoding

- ❑ Label Encoding
- ❑ One Hot Encoding

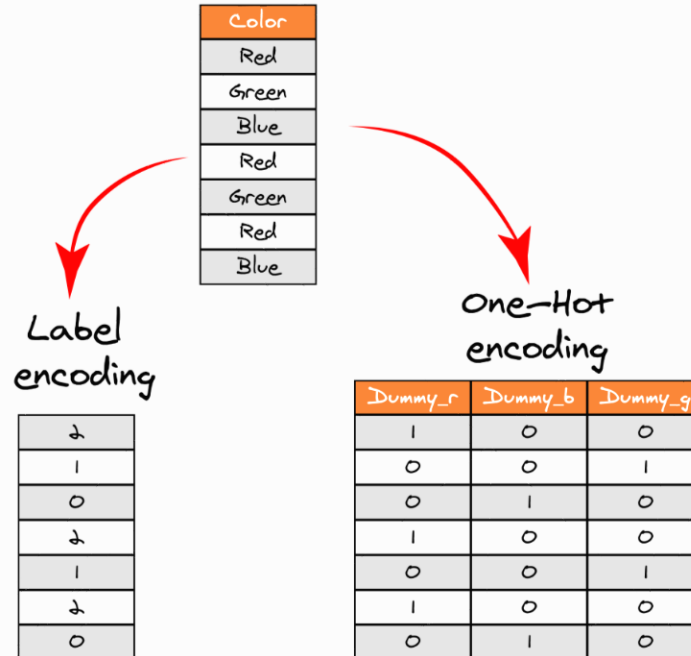
INTRODUCTION

Pre-Processing

Model Training

RESULTS

CONCLUSION



# Data Encoding using One Hot Encoding

INTRODUCTION

- ☐ Pclass Column
- ☐ Embarked Column
- ☐ Sex Column
- ☐ Title Column

Pre-Processing

Model Training

RESULTS

CONCLUSION

```
titanic_df_18 = titanic_df_17.copy()

# Select categorical columns
categorical_columns = ['Pclass', 'Sex', 'Embarked', 'Title']

# Extract categorical features
categorical_data = titanic_df_18[categorical_columns]

# Initialize OneHotEncoder
one_hot_encoder = OneHotEncoder(sparse=False, drop='first')

# Fit and transform the categorical features
encoded_data = one_hot_encoder.fit_transform(categorical_data)

# Get the feature names from OneHotEncoder
feature_names = one_hot_encoder.get_feature_names_out(categorical_columns)

# Create DataFrame with encoded features
encoded_df = pd.DataFrame(encoded_data, columns=feature_names, index=titanic_df_18.index)

# Drop original categorical columns from titanic_df_18
titanic_df_18.drop(columns=categorical_columns, inplace=True)

# Concatenate titanic_df_18 with encoded_df
titanic_df_18 = pd.concat([titanic_df_18, encoded_df], axis=1)
```



# Feature Selection

INTRODUCTION

- ❑ Correlation Matrix and Heat Map Visualization
- ❑ Mutual Information Feature Selection

Pre-Processing

Model Training

RESULTS

CONCLUSION

## Mutual Information

- Captures both linear and non-linear relationships between variables
- Considers the direct relationship between features and the target variable
- Computationally more expensive compared to correlation matrix
- May require tuning parameters like the number of features to select (k)

## Correlation Matrix


- Helps identify highly correlated features
- Easy to understand and visualize
- Only captures linear relationships between variables
- May not identify non-linear relationships

# Feature Selection

INTRODUCTION

## Correlation Matrix and Heat Map Visualization

Pre-Processing



Model Training

RESULTS

CONCLUSION

# Feature Selection

INTRODUCTION


## ❑ Mutual Information Feature Selection

Pre-Processing

Model Training

RESULTS

CONCLUSION



# Data Balancing

INTRODUCTION

Pre-Processing

Model Training

RESULTS

CONCLUSION

## Undersampling



Original

( 80 ) ( 45 )



Undersampling

( 45 ) ( 45 )

## Oversampling



Original

( 80 ) ( 45 )



Oversampling

( 80 ) ( 80 )

## SMOTE (Synthetic Minority Oversampling Technique)



Original

( 80 ) ( 45 )

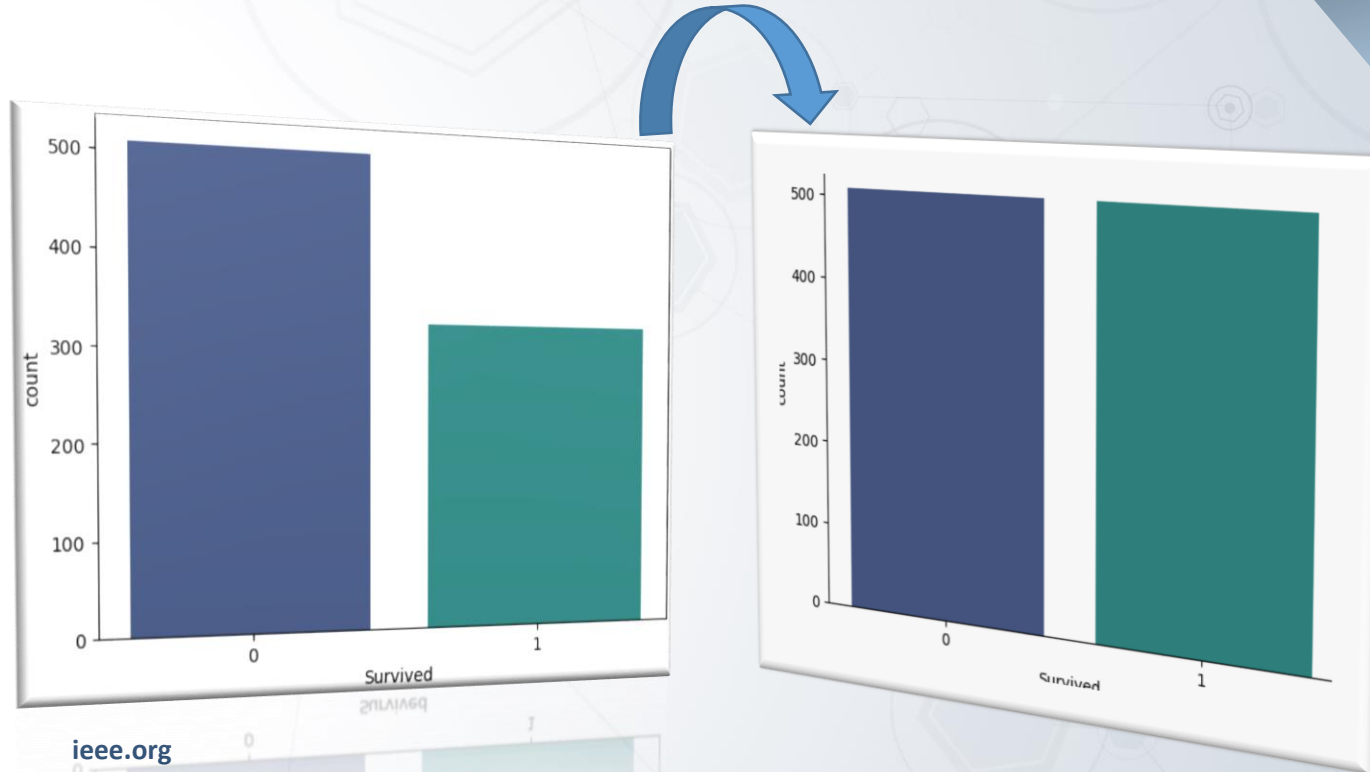


Oversampling

( 80 ) ( 80 )

# Data Balancing

## SMOTE (Synthetic Minority Oversampling Technique)



INTRODUCTION

Pre-Processing

Model Training

RESULTS

CONCLUSION

# Data Splitting



INTRODUCTION

**Pre-Processing**

Model Training

RESULTS

CONCLUSION



# Model Evaluation Technique


INTRODUCTION

Pre-Processing

**Model Training**

RESULTS

CONCLUSION



## K-Fold Cross-Validation

Cross-validation is a technique used in machine learning to check how well a model can generalize to new, or unseen data.



Testing Data

Training Data



# Model Evaluation Technique

## ❑ Add K-Fold Cross-Validation Technique for Better Accuracy

```
log_reg_model = LogisticRegression()  
num_splits = 3  
  
kf = KFold(n_splits=num_splits, shuffle=True, random_state=21)  
  
cv_scores = cross_val_score(log_reg_model, X, y, cv=kf).mean()  
  
print("Cross-validation scores:", cv_scores)
```

Cross-validation scores: 0.8128904249871992

INTRODUCTION

Pre-Processing

Model Training

RESULTS

CONCLUSION

# Modifications

- **Parameters:**

Kernal Function

Regularization Parameter

Tolerance to Convergence

Class Weight

Penalty

Solver

Multi-class Handling

Variable

Fixed

INTRODUCTION

Pre-Processing

**Model Training**

RESULTS

CONCLUSION

# Modifications

## ❑ Add Hyper Parameter Tuning

```
param_grid = {
    'penalty': ['l1', 'l2', 'elasticnet', 'none'],
    'C': np.logspace(-2, 2, 5),
    'solver': ['lbfgs', 'liblinear', 'saga'],
    'max_iter': [100, 500, 1000],
    'l1_ratio': np.linspace(0, 1, 3),
}

num_splits = 3
kf = KFold(n_splits=num_splits, shuffle=True, random_state=21)

grid_search = GridSearchCV(log_reg_model, param_grid, cv=kf, scoring='accuracy')
grid_search.fit(X_train, y_train)

best_params = grid_search.best_params_

cv_scores = grid_search.best_score_
best_model = grid_search.best_estimator_

print("Best Hyperparameters" )
best_params
```

INTRODUCTION

Pre-Processing

Model Training

RESULTS

CONCLUSION

Best Hyperparameters

```
{'C': 0.01,
 'l1_ratio': 0.0,
 'max_iter': 100,
 'penalty': 'none',
 'solver': 'lbfgs'}
```

# Performance Evaluation Metrics

INTRODUCTION

Pre-Processing

Model Training

RESULTS

CONCLUSION

☐ Accuracy    ☐ Precision    ☐ Recall    ☐ F1 - Score

True Class		Function Name	Formula
Predicted Class	True Positive (TP)	Accuracy	$\frac{TP + TN}{TP + TN + FP + FN}$
	False Positive (FP)	Precision	$\frac{TP}{TP + FP}$
	False Negative (FN)	Recall	$\frac{TP}{TP + FN}$
	True Negative (TN)	F1 score	$\frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$


# Key Results

## ❑ Before Hyper parameter Tuning

Training Accuracy: 0.8286  
Training Precision: 0.7968  
Training Recall: 0.7576  
Training F1-score: 0.7767  
Training ROC AUC Score: 0.8161

Testing Accuracy: 0.8274  
Testing Precision: 0.7937  
Testing Recall: 0.7576  
Testing F1-score: 0.7752  
Testing ROC AUC Score: 0.8151

Training vs Testing Metrics



INTRODUCTION

Pre-Processing

Model Training

## RESULTS

CONCLUSION

# Key Results

## ❑ After Hyper parameter Tuning

Training Accuracy after Hyperparameter Tuning: 0.8286  
Training Precision after Hyperparameter Tuning: 0.7945  
Training Recall after Hyperparameter Tuning: 0.7614  
Training F1-score after Hyperparameter Tuning: 0.7776  
Training ROC AUC Score after Hyperparameter Tuning: 0.8168

Testing Accuracy after Hyperparameter Tuning: 0.8333  
Testing Precision after Hyperparameter Tuning: 0.7969  
Testing Recall after Hyperparameter Tuning: 0.7727  
Testing F1-score after Hyperparameter Tuning: 0.7846  
Testing ROC AUC Score after Hyperparameter Tuning: 0.8226

INTRODUCTION

Pre-Processing

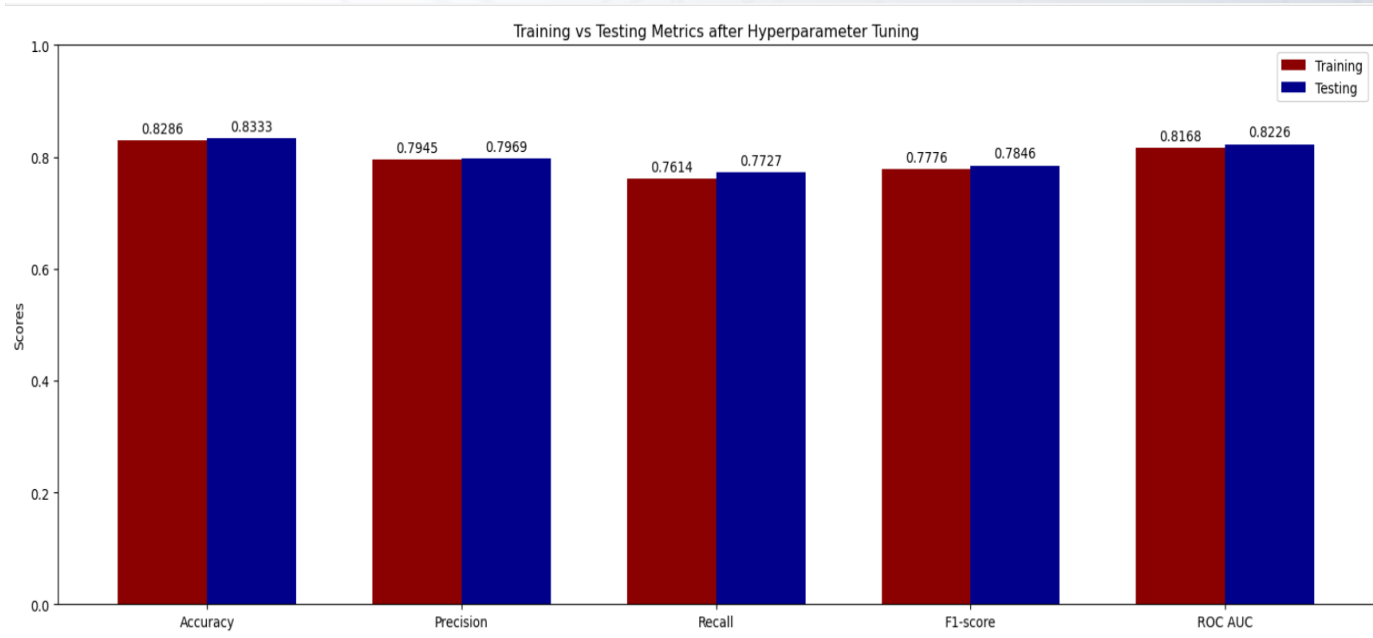
Model Training

**RESULTS**

CONCLUSION

# Key Results

## ❑ After Hyper parameter Tuning



INTRODUCTION

Pre-Processing

Model Training

RESULTS

CONCLUSION

# Thank You

INTRODUCTION

METHODOLOGY

RESULTS

EPICS

CONCLUSION